

# ----- Training -----

NOTE: This section is for practising. Questions specified by Q are not a part of your coursework. You can check the comments to see what each cell does.

## Importing necessary Libraries

```
In [1]: import pandas as pd
import numpy as np
```

## Python Warmup exercises

- Q1: Write a code that receives an optional text, print all letters in the text and store them in a list

```
In [2]: output=[]
text1='''welcome to workshop 1'''
for i in text1:
    print(i)
    output.append(i)
print(output)
```

w  
e  
l  
c  
o  
m  
e

t  
o

w  
o  
r  
k  
s  
h  
o  
p

1  
['w', 'e', 'l', 'c', 'o', 'm', 'e', ' ', 't', 'o', ' ', 'w', 'o', 'r', 'k',  
's', 'h', 'o', 'p', ' ', '1']

- Q2: Write a code that receives an optional text, prints all words in the text and store them in a list

```
In [3]: output=[]
text1='''welcome to workshop 1'''
for word in text1.split():
    print(word)
    output.append(word)
print(output)

welcome
to
workshop
1
['welcome', 'to', 'workshop', '1']
```

- Q3: Write Q2 answer as a function

```
In [4]: text1='''welcome to workshop 1'''
def print_word(text):
    for word in text.split():
        print(word)
print_word(text1)

welcome
to
workshop
1
```

- Q4: Write a program to find out the average of a set of integers. The program prompts the user to enter the count of numbers

```
In [6]: count = int(input("Enter the count of numbers: ")) # This gets the count of numbers
# puts it in an integer variable
i = 0
sum = 0
for i in range(count):
    x = int(input("Enter an integer: "))
    sum = sum + x
avg = sum/count
print("The average is: ", avg)

Enter the count of numbers: 1
Enter an integer: 1
The average is: 1.0
```

- Q5: Write a program to check whether the given number is even or not.

```
In [7]: # If we divide an even number by 2, the remainder will be zero
number = input("Enter a number: ")
x = int(number)%2
if x == 0:
    print("The number is Even.")
else:
    print("The number is Odd.")
```

Enter a number: 1  
The number is Odd.

- Q6: Write a program to find the average of 4 numbers using while loop. The program prompts the user to enter the numbers

```
In [8]: count = 0
sum = 0.0
while(count<4):
    number = float(input("Enter a real number: "))
    count=count+1
    sum = sum+number
avg = sum/10
print("Average is :",avg)
```

Enter a real number: 1  
Enter a real number: 1  
Enter a real number: 1  
Enter a real number: 1  
Average is : 0.4

- Q7: Write a function to find the average of a list of integers

```
In [9]: def give_me_average(in_list):
    count=len(in_list)
    sum=0
    for i in in_list:
        sum=sum+i
    avg = sum/count
    return avg
a=[1,2,3,4]
print("Average is :",give_me_average(a))
```

Average is : 2.5

- Q8: Write a program to display all integers within the range 100-120 whose sum of digits is an even number and store them in a list

```
In [10]: output=[]
for i in range(100,120):
    num = i
    sum = 0
    while(num!=0):
        digit = num%10
        sum = sum + digit
        num = num//10
    if(sum%2==0):
        output.append(i)
        print(i)
print(output)
```

```
101
103
105
107
109
110
112
114
116
118
[101, 103, 105, 107, 109, 110, 112, 114, 116, 118]
```

- Q9: Write a program to check whether the given integer is a multiple of both 5 and 7

```
In [12]: number = int(input("Enter an integer: "))
if((number%5==0)and(number%7==0)):
    print(number, "is a multiple of both 5 and 7")
else:
    print(number, "is not a multiple of both 5 and 7")
```

```
Enter an integer: 1
1 is not a multiple of both 5 and 7
```

## Pandas warmup exercises

- You can install Pandas on your system by using pip as follows:  
pip install pandas

- You can import Pandas as follows:  
import pandas

- Pandas is usually imported under the pd alias.  
import pandas as pd

- You can check Pandas version as follows:  
import pandas as pd  
print(pd.\_\_version\_\_)

```
In [13]: # You can use the DataFrame() method to create a Pandas DataFrame from a dictio
data = {"calories": [420, 380, 390], "duration": [50, 40, 45]}
#Load data into a DataFrame object:
df = pd.DataFrame(data)
print(df)
print('\n') # Addine an empty line to the end of the data frame
# ---
data = [[420, 380, 390],[50, 40, 45]]
df = pd.DataFrame(data)
print(df)
```

	calories	duration
0	420	50
1	380	40
2	390	45

	0	1	2
0	420	380	390
1	50	40	45

```
In [14]: # You can use loc() or iloc() to return values, rows, or columns inside a data
# column name, but using iloc() you only need to specify the column index.
data = {"calories": [420, 380, 390], "duration": [50, 40, 45]}
#Load data into a DataFrame object:
df = pd.DataFrame(data)
print(df)
print("\n")
print(df.loc[0,'calories'])
print("\n")
print(df.loc[0,:])
print("\n")
print(df.iloc[:,0])
```

	calories	duration
0	420	50
1	380	40
2	390	45

420

calories 420  
duration 50  
Name: 0, dtype: int64

0 420  
1 380  
2 390  
Name: calories, dtype: int64

In [15]: *# You can get the indices and columns names using index and column attributes a*

```
data = {"calories": [420, 380, 390], "duration": [50, 40, 45]}
#Load data into a DataFrame object:
df = pd.DataFrame(data)
index_names=list(df.index)
column_names=list(df.columns)
print(index_names,end='\n\n')
print(column_names)
```

```
[0, 1, 2]
```

```
['calories', 'duration']
```

In [16]: *# With the index and columns argument, you can name your own indexes and column*

```
data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
df = pd.DataFrame(data, index = ["day1", "day2", "day3"])
print(df)
```

	calories	duration
day1	420	50
day2	380	40
day3	390	45

In [17]: *# With the index argument, you can name your own indexes and columns*

```
data = [[420, 380, 390], [50, 40, 45]]
df = pd.DataFrame(data, columns = ["c1", "c2", "c3"])
print(df)
```

	c1	c2	c3
0	420	380	390
1	50	40	45

## Importing a dataset for this workshop

### Importing the dataset directly from the source

In [18]: *# Installing a package which is needed to download the dataset from its online*  
*# URL : <https://archive.ics.uci.edu/dataset/2/adult>*  
*!pip3 install ucimlrepo*

Requirement already satisfied: ucimlrepo in c:\users\hp\anaconda3\lib\site-packages (0.0.3)

```
In [19]: # Downloading the dataset from the online source. The first two lines are given
from ucimlrepo import fetch_ucirepo
# fetch dataset
adult = fetch_ucirepo(id=2)

# Putting data in a pandas dataframe
X = adult.data.features
y = adult.data.targets
data=pd.concat([X,y],axis=1)

data
```

Out[19]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...	...	...	...	...	...
48837	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
48838	64	NaN	321403	HS-grad	9	Widowed	NaN	Other-relative	Black
48839	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
48840	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
48841	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

48842 rows × 15 columns

In [20]: `#printing data`  
`data`

Out[20]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...	...	...	...	...	...
48837	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
48838	64	NaN	321403	HS-grad	9	Widowed	NaN	Other-relative	Black
48839	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
48840	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
48841	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

48842 rows × 15 columns



## Importing the dataset from a folder on your local disk



In [21]: `#importing the dataset as a Pandas DataFrame into Python if the dataset is stored locally`  
`# You can download the dataset from the following URL: https://archive.ics.uci.edu/dataset/1/adult`  
`data = pd.read_csv('C:/Users/HP/Downloads/adult.csv') # Replace the current path with the local path to the dataset`



## Exploring the dataset

```
In [22]: # Showing the first 5 rows of the dataset  
data.head()
```

Out[22]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female

```
In [23]: # Finding the shape of the data  
print(data.shape)
```

(48842, 15)

```
In [24]: # Generate a dataset by randomly extracting 30000 rows (samples)  
data_new = data.sample(n=30000, random_state = 48)
```

In [25]: *# Printing the new dataset*  
data\_new

Out[25]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	
<b>8029</b>	29	Private	216481	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	I
<b>45203</b>	36	Private	280570	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	
<b>27498</b>	25	?	100903	Bachelors	13	Married-civ-spouse	?	Wife	White	I
<b>48416</b>	47	Private	145636	Assoc-voc	11	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>43230</b>	33	Private	119422	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	
...	...	...	...	...	...	...	...	...	...	...
<b>37544</b>	20	Private	166371	HS-grad	9	Never-married	Craft-repair	Other-relative	White	
<b>20482</b>	80	Private	202483	HS-grad	9	Married-spouse-absent	Adm-clerical	Not-in-family	White	I
<b>39667</b>	20	Private	175808	HS-grad	9	Never-married	Craft-repair	Own-child	White	
<b>40001</b>	25	State-gov	31350	Some-college	10	Never-married	Other-service	Not-in-family	White	
<b>48523</b>	19	Private	239057	HS-grad	9	Never-married	Craft-repair	Own-child	White	

30000 rows × 15 columns



In [26]: *# The indices of different rows in the dataset are currently messy. This happens if the dataset is not sorted by index.*  
*# the dataset if you are unsure the indices are correct.*  
data\_new.reset\_index(drop=True, inplace=True)

In [27]: *# Checking if the indices are correct*  
data\_new

Out[27]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	
<b>0</b>	29	Private	216481	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	I
<b>1</b>	36	Private	280570	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	
<b>2</b>	25	?	100903	Bachelors	13	Married-civ-spouse	?	Wife	White	I
<b>3</b>	47	Private	145636	Assoc-voc	11	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>4</b>	33	Private	119422	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	
...	...	...	...	...	...	...	...	...	...	
<b>29995</b>	20	Private	166371	HS-grad	9	Never-married	Craft-repair	Other-relative	White	
<b>29996</b>	80	Private	202483	HS-grad	9	Married-spouse-absent	Adm-clerical	Not-in-family	White	I
<b>29997</b>	20	Private	175808	HS-grad	9	Never-married	Craft-repair	Own-child	White	
<b>29998</b>	25	State-gov	31350	Some-college	10	Never-married	Other-service	Not-in-family	White	
<b>29999</b>	19	Private	239057	HS-grad	9	Never-married	Craft-repair	Own-child	White	

30000 rows × 15 columns



In [28]: data

Out[28]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...	...	...	...	...	...
48837	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
48838	64	NaN	321403	HS-grad	9	Widowed	NaN	Other-relative	Black
48839	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
48840	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
48841	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

48842 rows × 15 columns



In [29]: *# Getting statistical information of the dataset for different columns (feature data.describe(include="all"))*

Out[29]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	re
<b>count</b>	48842.000000	47879	4.884200e+04	48842	48842.000000	48842	47876	
<b>unique</b>	NaN	9	NaN	16	NaN	7	15	
<b>top</b>	NaN	Private	NaN	HS-grad	NaN	Married-civ-spouse	Prof-specialty	
<b>freq</b>	NaN	33906	NaN	15784	NaN	22379	6172	
<b>mean</b>	38.643585	NaN	1.896641e+05	NaN	10.078089	NaN	NaN	
<b>std</b>	13.710510	NaN	1.056040e+05	NaN	2.570973	NaN	NaN	
<b>min</b>	17.000000	NaN	1.228500e+04	NaN	1.000000	NaN	NaN	
<b>25%</b>	28.000000	NaN	1.175505e+05	NaN	9.000000	NaN	NaN	
<b>50%</b>	37.000000	NaN	1.781445e+05	NaN	10.000000	NaN	NaN	
<b>75%</b>	48.000000	NaN	2.376420e+05	NaN	12.000000	NaN	NaN	
<b>max</b>	90.000000	NaN	1.490400e+06	NaN	16.000000	NaN	NaN	

In [30]: *# Dropping row with Null values (This topic will be discussed in detail in the n*  
data.dropna(inplace=True)

In [31]: data

Out[31]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...	...	...	...	...	...
48836	33	Private	245211	Bachelors	13	Never-married	Prof-specialty	Own-child	White
48837	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
48839	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
48840	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
48841	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

47621 rows × 15 columns



NOTE: As you can see, because a few rows have been dropped, the index numbers don't correspond to the actual row number (the number of rows is 47623 but the last index is 48841). We need to reindex the data

```
In [32]: # reindexing the original data
data.reset_index(drop=True, inplace=True)
```

In [33]: data

Out[33]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
<b>0</b>	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
<b>1</b>	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
<b>2</b>	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
<b>3</b>	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
<b>4</b>	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...	...	...	...	...	...
<b>47616</b>	33	Private	245211	Bachelors	13	Never-married	Prof-specialty	Own-child	White
<b>47617</b>	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
<b>47618</b>	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
<b>47619</b>	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
<b>47620</b>	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

47621 rows × 15 columns



```
In [34]: # Getting the number of Null values for different columns
data.isna().sum()
```

```
Out[34]: age                0
workclass                0
fnlwgt                  0
education               0
education-num           0
marital-status          0
occupation              0
relationship            0
race                   0
sex                    0
capital-gain            0
capital-loss            0
hours-per-week          0
native-country          0
income                 0
dtype: int64
```

```
In [35]: # Showing the dataset information
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47621 entries, 0 to 47620
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   47621 non-null  int64
 1   workclass              47621 non-null  object
 2   fnlwgt                 47621 non-null  int64
 3   education              47621 non-null  object
 4   education-num          47621 non-null  int64
 5   marital-status         47621 non-null  object
 6   occupation              47621 non-null  object
 7   relationship           47621 non-null  object
 8   race                   47621 non-null  object
 9   sex                    47621 non-null  object
10   capital-gain           47621 non-null  int64
11   capital-loss           47621 non-null  int64
12   hours-per-week         47621 non-null  int64
13   native-country         47621 non-null  object
14   income                 47621 non-null  object
dtypes: int64(6), object(9)
memory usage: 5.4+ MB
```

```
In [36]: data.shape
```

```
Out[36]: (47621, 15)
```



```
In [37]: # Getting the count of different values in the column "education-num"
data['education-num'].value_counts()
```

```
Out[37]: education-num
9      15444
10     10512
13      7881
14      2610
11      2034
7       1746
12      1566
6       1336
4        912
15       819
5        735
8        633
16       582
3        494
2        239
1         78
Name: count, dtype: int64
```

```
In [38]: # Getting the count of different values in the column "education"
data['education'].value_counts()
```

```
Out[38]: education
HS-grad      15444
Some-college 10512
Bachelors    7881
Masters      2610
Assoc-voc    2034
11th         1746
Assoc-acdm   1566
10th         1336
7th-8th      912
Prof-school  819
9th          735
12th         633
Doctorate    582
5th-6th      494
1st-4th      239
Preschool    78
Name: count, dtype: int64
```

```
In [39]: # Dropping a column
data = data.drop(['fnlwtg'], axis=1)
```

```
In [40]: data.shape
```

```
Out[40]: (47621, 14)
```

```
In [41]: # Getting the number of unique values of a column  
data['education'].nunique()
```

Out[41]: 16

```
In [42]: # Finding how many rows are related to either gender  
data['sex'].value_counts()
```

Out[42]: sex  
Male 31937  
Female 15684  
Name: count, dtype: int64

```
In [43]: # Calculating the average age of different genders in the dataset  
data['age'].groupby([data['sex']]).mean()
```

Out[43]: sex  
Female 36.961043  
Male 39.465542  
Name: age, dtype: float64

```
In [44]: # Getting the average age of different genders in the dataset broken down based
data['age'].groupby([data['sex'],data['education']]).mean()
```

```
Out[44]: sex      education
Female 10th      36.520642
       11th      29.891410
       12th      29.716418
       1st-4th   46.508475
       5th-6th   45.346774
       7th-8th   51.088889
       9th       41.566038
       Assoc-acdm 36.402619
       Assoc-voc  38.000000
       Bachelors  35.747837
       Doctorate  45.557522
       HS-grad    38.862285
       Masters    42.520482
       Preschool  44.409091
       Prof-school 40.131783
       Some-college 33.778719
Male   10th      38.406667
       11th      33.479185
       12th      33.842593
       1st-4th   46.050000
       5th-6th   42.524324
       7th-8th   48.608443
       9th       40.556405
       Assoc-acdm 38.537173
       Assoc-voc  38.853881
       Bachelors  40.190319
       Doctorate  47.936034
       HS-grad    39.073626
       Masters    44.584270
       Preschool  41.160714
       Prof-school 46.017391
       Some-college 37.176643
Name: age, dtype: float64
```

```
In [45]: # Getting the maximum age of different races in the dataset
data['age'].groupby([data['race']]).max()
```

```
Out[45]: race
Amer-Indian-Eskimo    82
Asian-Pac-Islander    90
Black                 90
Other                 77
White                 90
Name: age, dtype: int64
```

```
In [46]: # Extracting the age and education columns and creating a new DataFrame using t
a=data['age']
b=data['education']
new_data=pd.concat([a,b],axis=1)
```

In [47]: data

Out[47]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female
...	...	...	...	...	...	...	...	...	...
47616	33	Private	Bachelors	13	Never-married	Prof-specialty	Own-child	White	Male
47617	39	Private	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White	Female
47618	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White	Male
47619	44	Private	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander	Male
47620	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male

47621 rows × 14 columns



- Q10: Write a function that receives the dataset and replace Famle with F and Male with M (please try to write it yourself before checking the answer in the next cell)

```
In [48]: # Answer
def encode_sex(data):
    rows=data.shape[0]
    a=data.loc[:, 'sex']
    for i in range(rows):
        if a[i]=="Male":
            data.loc[i, "sex"]="M"
        elif a[i]=="Female":
            data.loc[i, "sex"]="F"
    return data
```

```
In [49]: # Copying the data
data_copy=data.copy()
```

```
In [50]: # Applying the encode_sex function to the copied data
data_encoded=encode_sex(data_copy)
data_encoded.head()
```

Out[50]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	M	2174
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	M	0
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	M	0
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	M	0
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	F	0

In [51]: data

Out[51]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female
...	...	...	...	...	...	...	...	...	...
47616	33	Private	Bachelors	13	Never-married	Prof-specialty	Own-child	White	Male
47617	39	Private	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White	Female
47618	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White	Male
47619	44	Private	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander	Male
47620	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male

47621 rows × 14 columns



# Tasks

## TASK 1.1: Import the dataset from the URL we used in this workshop. Then generate a new dataset by randomly extracting 10000 samples. Reindex the generated dataset and remove NULL values. Name the new dataset 'task\_dataset'

```
In [52]: ##### WRITE THE CODE IN THIS CELL #####  
# Downloading the dataset from the online source. The first two lines are given  
from ucimlrepo import fetch_ucirepo  
# fetch dataset  
adult = fetch_ucirepo(id=2)  
  
# Putting data in a pandas dataframe  
X = adult.data.features  
y = adult.data.targets  
data=pd.concat([X,y],axis=1)  
  
#Printing data  
data
```

Out[52]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...	...	...	...	...	...
48837	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
48838	64	NaN	321403	HS-grad	9	Widowed	NaN	Other-relative	Black
48839	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
48840	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
48841	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

48842 rows × 15 columns

##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####

The adult dataset was fetched from the UCI Machine learning repository ucimlrepo, and i loaded it into a pandas dataframe, and splitted to set all the features columns as X and the target column (Y) as income. i concatenated (X,y) back into a single dataframe and then stored into Data, which was printed and shown above as containing 48842 rows and 15 columns.

i went further to generate new dataset by randomly extracting 10000 rows using the sample function. The data was then renamed, stored into data\_new and printed. The randomly sampled data (data\_new) was reindexed.

further, data preprocessing and cleaning was carried out by dropping all null values in the dataset and the cleaned data stored as task\_dataset



(120 WORDS)

```
In [53]: # Generate a dataset by randomly extracting 10000 rows (samples)
data_new = data.sample(n=10000, random_state = 48)

# Printing the new dataset
data_new
```

Out[53]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
<b>8029</b>	29	Private	216481	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White
<b>45203</b>	36	Private	280570	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White
<b>27498</b>	25	?	100903	Bachelors	13	Married-civ-spouse	?	Wife	White
<b>48416</b>	47	Private	145636	Assoc-voc	11	Married-civ-spouse	Handlers-cleaners	Husband	White
<b>43230</b>	33	Private	119422	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White
...	...	...	...	...	...	...	...	...	...
<b>5637</b>	19	Private	63363	Some-college	10	Never-married	Sales	Own-child	White
<b>46277</b>	53	Private	58535	HS-grad	9	Divorced	Sales	Not-in-family	White
<b>26200</b>	30	Private	342709	HS-grad	9	Married-civ-spouse	Handlers-cleaners	Husband	White
<b>33519</b>	41	Self-emp-not-inc	134724	Assoc-voc	11	Married-civ-spouse	Other-service	Wife	White
<b>28113</b>	21	Private	252253	Some-college	10	Never-married	Adm-clerical	Unmarried	Black

10000 rows × 15 columns



```
In [54]: # Reindexing the indices and printing the outcome
data_new.reset_index(drop=True, inplace=True)

#printing outcome
data_new
```

Out[54]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	
<b>0</b>	29	Private	216481	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	F
<b>1</b>	36	Private	280570	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	
<b>2</b>	25	?	100903	Bachelors	13	Married-civ-spouse	?	Wife	White	F
<b>3</b>	47	Private	145636	Assoc-voc	11	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>4</b>	33	Private	119422	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	
...	...	...	...	...	...	...	...	...	...	
<b>9995</b>	19	Private	63363	Some-college	10	Never-married	Sales	Own-child	White	F
<b>9996</b>	53	Private	58535	HS-grad	9	Divorced	Sales	Not-in-family	White	F
<b>9997</b>	30	Private	342709	HS-grad	9	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>9998</b>	41	Self-emp-not-inc	134724	Assoc-voc	11	Married-civ-spouse	Other-service	Wife	White	F
<b>9999</b>	21	Private	252253	Some-college	10	Never-married	Adm-clerical	Unmarried	Black	F

10000 rows × 15 columns



```
In [55]: # Dropping row with Null values and renaming
data_new.dropna(inplace=True)
```

```
In [56]: # Renaming the dataset
task_dataset = data_new

#printing
task_dataset
```

Out[56]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	
<b>0</b>	29	Private	216481	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	F
<b>1</b>	36	Private	280570	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	
<b>2</b>	25	?	100903	Bachelors	13	Married-civ-spouse	?	Wife	White	F
<b>3</b>	47	Private	145636	Assoc-voc	11	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>4</b>	33	Private	119422	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	
...	...	...	...	...	...	...	...	...	...	
<b>9995</b>	19	Private	63363	Some-college	10	Never-married	Sales	Own-child	White	F
<b>9996</b>	53	Private	58535	HS-grad	9	Divorced	Sales	Not-in-family	White	F
<b>9997</b>	30	Private	342709	HS-grad	9	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>9998</b>	41	Self-emp-not-inc	134724	Assoc-voc	11	Married-civ-spouse	Other-service	Wife	White	F
<b>9999</b>	21	Private	252253	Some-college	10	Never-married	Adm-clerical	Unmarried	Black	F

9765 rows × 15 columns



```
In [57]: # Reindexing the indices and printing the outcome
task_dataset.reset_index(drop=True, inplace=True)

#printing outcome
task_dataset
```

Out[57]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	
<b>0</b>	29	Private	216481	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	F
<b>1</b>	36	Private	280570	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	
<b>2</b>	25	?	100903	Bachelors	13	Married-civ-spouse	?	Wife	White	F
<b>3</b>	47	Private	145636	Assoc-voc	11	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>4</b>	33	Private	119422	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	
...	...	...	...	...	...	...	...	...	...	
<b>9760</b>	19	Private	63363	Some-college	10	Never-married	Sales	Own-child	White	F
<b>9761</b>	53	Private	58535	HS-grad	9	Divorced	Sales	Not-in-family	White	F
<b>9762</b>	30	Private	342709	HS-grad	9	Married-civ-spouse	Handlers-cleaners	Husband	White	
<b>9763</b>	41	Self-emp-not-inc	134724	Assoc-voc	11	Married-civ-spouse	Other-service	Wife	White	F
<b>9764</b>	21	Private	252253	Some-college	10	Never-married	Adm-clerical	Unmarried	Black	F

9765 rows × 15 columns



```
In [58]: # Getting the number of Null values for different columns  
task_dataset.isna().sum()
```

```
Out[58]: age                0  
workclass              0  
fnlwgt                 0  
education              0  
education-num          0  
marital-status         0  
occupation             0  
relationship           0  
race                   0  
sex                    0  
capital-gain           0  
capital-loss           0  
hours-per-week         0  
native-country         0  
income                 0  
dtype: int64
```

## TASK 1.2: Write a code to find how much contribution each sex and occupation category made to the capital-gain on average. Apply the code to the task\_dataset and report the result (Hint: you need to use the groupby method)

In [59]: `##### WRITE THE CODE IN THIS CELL #####`  
`# Getting the average capital gain based on the contribution from each sex and`  
`task_dataset['capital-gain'].groupby([task_dataset['sex'],task_dataset['occupat`

Out[59]:

sex	occupation	
Female	?	359.855556
	Adm-clerical	508.217678
	Craft-repair	273.969231
	Exec-managerial	1311.540984
	Farming-fishing	1049.523810
	Handlers-cleaners	199.136364
	Machine-op-inspct	290.217687
	Other-service	191.525394
	Priv-house-serv	74.425926
	Prof-specialty	1008.057269
	Protective-serv	5081.000000
	Sales	276.654822
	Tech-support	382.803419
	Transport-moving	533.800000
Male	?	879.403509
	Adm-clerical	425.640751
	Armed-Forces	0.000000
	Craft-repair	906.637306
	Exec-managerial	2339.700234
	Farming-fishing	441.013746
	Handlers-cleaners	149.265957
	Machine-op-inspct	577.642082
	Other-service	155.411765
	Priv-house-serv	118.800000
	Prof-specialty	3918.730198
	Protective-serv	368.563158
	Sales	1129.695105
	Tech-support	419.494565
	Transport-moving	409.625000

Name: capital-gain, dtype: float64

##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####

The result of the average capital gain of each sex based on occupation shows some economic disparities between the males and the females within same occupation. Occupations such as Admin clerical, farming-fishing, Handlers-cleaners, Protective serv and transport-moving, recorded higher average contribution from the females to the capital gain, than the males. on the average females in Admin Clerical made 12% more contribution than the males. 40% more in farming-fishing occupation, 14% more in Handlers-cleaners, 86% more in Protective-serv occupation and 14% more in Transport-moving.

While in occupations such as Craft repair, Exec-managerial, Machine-op-inspct, Priv-house-serv, Prof-specialty, Sales and Tech-support, the males recorded higher average contribution to the capital gain than the females with 54%, 28%, 33%, 24%, 60%, 60% and 4% respectively

(121 WORDS)

**TASK 1.3: Write a code to find the country with the highest number of people with a Bachelors degree. Apply the code to the task\_dataset and report the result (Hint: you need to use the groupby method)**

```
In [60]: ##### WRITE THE CODE IN THIS CELL #####  
# adjust column names based on the actual dataset  
bachelors_task_dataset = task_dataset[task_dataset['education'] == 'Bachelors']  
  
# Group by 'native-country' and count the number of people with a Bachelor's de  
country_bachelors_count = bachelors_task_dataset.groupby('native-country').size  
  
# Find the country with the highest count  
highest_country = country_bachelors_count.idxmax()  
highest_count = country_bachelors_count.max()  
  
# Sort the counts in descending order  
country_bachelors_count_sorted = country_bachelors_count.sort_values(ascending=  
  
# Print the sorted counts  
print(country_bachelors_count_sorted)  
  
print(f"The country with the highest number of people with a Bachelors degree i
```



native-country	
United-States	1456
?	30
Philippines	15
England	9
Japan	8
Canada	7
Mexico	7
India	7
Italy	6
China	6
South	6
Puerto-Rico	6
Cuba	5
Taiwan	5
Germany	5
Iran	4
El-Salvador	4
Jamaica	3
Vietnam	3
Cambodia	2
Peru	2
Dominican-Republic	2
Thailand	2
Ireland	1
Guatemala	1
Greece	1
Poland	1
Ecuador	1
Columbia	1

dtype: int64

The country with the highest number of people with a Bachelors degree is United-States with 1456 people.

##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####

The codes above sought to find the country with the highest distribution of individuals with a Bachelor's degree across different countries from the task\_dataset

This dataset is converted into a pandas DataFrame, task\_dataset, which represents the larger dataset containing various education levels and countries of origin for different individuals. The dataset is then filtered to select only the rows where the education column has the value 'Bachelors'. This filtered dataset is stored in bachelors\_task\_dataset.

The code then grouped the filtered dataset by the native-country column, which categorizes individuals based on their country of origin. The size() function is used to count the number of individuals with a Bachelor's degree in each country, resulting in a Series object, country\_bachelors\_count, indexed by country name with counts as values.

To find the country with the highest number of individuals holding a Bachelor's degree, i used the `idxmax()` on `country_bachelors_count`. This function returns the country name of the first occurrence of the maximum value followed by the highest count of individuals with a Bachelor's degree which is obtained using `max()`.

Finally, The counts of individuals with a Bachelor's degree by country were sorted in descending order using `sort_values(ascending=False)`. This sorted Series, `country_bachelors_count_sorted`, simplifies the distribution across different countries, and the United States was identified as the country with the highest count of individuals with Bachelors degree.

(222 WORDS)

## TASK 1.4: Write a code to receive two lists including five names and their respective ages and print 'Hello Name Age'

For example, if it received a list of two names ['Amin', 'Michael'] and respective ages [27,38], it would print 'Hello Amin 27', 'Hello Michael 38'. Each hello statement should be printed in a new line

In [61]: `##### WRITE THE CODE IN THIS CELL #####`

```
# creating two lists containing names and ages
names = ['Chienyeze', 'Ada', 'Ogechukwu', 'Chioma', 'Emeka']
ages = [38, 33, 39, 34, 40]

# Iterating through the lists and printing 'Hello Name Age'
for name, age in zip(names, ages):
    print(f'Hello {name} {age}')
```

```
Hello Chienyeze 38
Hello Ada 33
Hello Ogechukwu 39
Hello Chioma 34
Hello Emeka 40
```

`##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####`

The code received two lists including five names and their ages, and iterated through them using the zip method to print "Hello, 'Name', 'Age'"

## TASK 1.5: Write a code to receive an optional text, capitalise all words in the text and print them

```
In [62]: ##### WRITE THE CODE IN THIS CELL #####  
#defining the optional text parameters for words capitalisation  
def capitalise_words(text=None):  
    if text is None:  
        print("No text.")  
    elif text.isalpha():  
        capitalised_text = text.upper()  
        print("Capitalised Words:", capitalised_text)  
    else:  
        print("numbers cannot be capitalised")  
  
text = input("Enter optional text: ")  
capitalise_words(text)
```

```
Enter optional text: chienyeze  
Capitalised Words: CHIENYEZE
```

```
##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####
```

The code defines a function `capitalise_words` that takes an optional parameter `text` and performs capitalisation (`text.upper`) operations and others based on the content of `text`. if `text` is alpha, it capitalises it, if not alpha, it return as 'numbers cannot be capitalised' and if no text input, it returns as 'No text'. Using the `def`, `if`, `elif` and `else` functions, this optional text code was executed

(66 WORDS)

## TASK 1.6: Write a function to split the task\_dataset in half column-wise and swap the first half and the second half

```
In [63]: ##### WRITE THE CODE IN THIS CELL #####  
#splitting task_dataset into 2 halves and then swapping both halves  
def split_column(task_dataset):  
    #getthing the half  
    half_value = int(len(task_dataset.columns) // 2)  
    print(half_value)  
    #splitting the halves  
    print(task_dataset.iloc[:, half_value:].join(task_dataset.iloc[:, :half_val  
  
split_column(task_dataset)
```

```

7
      relationship  race    sex  capital-gain  capital-loss  \
0            Wife  White  Female           0           0
1        Husband  White   Male           0           0
2            Wife  White  Female           0           0
3        Husband  White   Male           0           0
4        Husband  White   Male           0           0
...           ...    ...    ...           ...           ...
9760    Own-child  White  Female           0           0
9761  Not-in-family  White  Female           0           0
9762        Husband  White   Male           0           0
9763            Wife  White  Female        3103           0
9764    Unmarried  Black  Female           0           0

      hours-per-week  native-country  income  age      workclass  fnlwgt  \
0                40  United-States  >50K   29      Private  216481
1                45  United-States  <=50K   36      Private  280570
2                25  United-States  <=50K   25           ?  100903
3                48  United-States  >50K   47      Private  145636
4                40  United-States  <=50K   33      Private  119422
...           ...    ...    ...    ...    ...    ...
9760            30  United-States  <=50K   19      Private   63363
9761            40  United-States  <=50K   53      Private   58535
9762            40  United-States  <=50K   30      Private  342709
9763            40  United-States  >50K   41  Self-emp-not-inc 134724
9764            40  United-States  <=50K   21      Private  252253

      education  education-num  marital-status      occupation
0        Masters             14  Married-civ-spouse  Exec-managerial
1    Some-college             10  Married-civ-spouse  Craft-repair
2        Bachelors             13  Married-civ-spouse           ?
3        Assoc-voc             11  Married-civ-spouse  Handlers-cleaners
4          HS-grad              9  Married-civ-spouse  Exec-managerial
...           ...    ...    ...    ...
9760    Some-college             10    Never-married      Sales
9761          HS-grad              9      Divorced      Sales
9762          HS-grad              9  Married-civ-spouse  Handlers-cleaners
9763        Assoc-voc             11  Married-civ-spouse  Other-service
9764    Some-college             10    Never-married  Adm-clerical

```

[9765 rows x 15 columns]

##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####

The swap column function was used to swap columns for the task\_dataset after splitting it into 2 halves columns-wise. the swapped halves are the concatenated along the column axis to form the final swap dataset (split\_column). the function enable the rearrangement of column data as needed

(46 WORDS)

**Task 1.7: Write a function that receives two numerical columns' names and compare their values for all rows. If the value of the first column is greater than the second column, the function should produce True, otherwise, it should produce False. The function should append an additional column to the dataset to store the results of the comparison for all rows. Apply the function to the "age" and "hours-per-week" columns in the task\_dataset and print the result.**

```
In [64]: ##### WRITE THE CODE IN THIS CELL#####  
# Comparing 'age' and 'hours-per-week' columns  
def compare_age_hours(col_a, col_b):  
    task_dataset['age_vs_hours'] = col_a > col_b  
    return task_dataset  
  
# Calling the function to compare 'age' and 'hours-per-week' columns  
result_task_dataset = compare_age_hours(task_dataset['age'], task_dataset['hours-per-week'])  
  
# Printing the result  
print(result_task_dataset)
```

	age	workclass	fnlwgt	education	education-num	\
0	29	Private	216481	Masters	14	
1	36	Private	280570	Some-college	10	
2	25	?	100903	Bachelors	13	
3	47	Private	145636	Assoc-voc	11	
4	33	Private	119422	HS-grad	9	
...	...	...	...	...	...	
9760	19	Private	63363	Some-college	10	
9761	53	Private	58535	HS-grad	9	
9762	30	Private	342709	HS-grad	9	
9763	41	Self-emp-not-inc	134724	Assoc-voc	11	
9764	21	Private	252253	Some-college	10	

	marital-status	occupation	relationship	race	sex	\
0	Married-civ-spouse	Exec-managerial	Wife	White	Female	
1	Married-civ-spouse	Craft-repair	Husband	White	Male	
2	Married-civ-spouse	?	Wife	White	Female	
3	Married-civ-spouse	Handlers-cleaners	Husband	White	Male	
4	Married-civ-spouse	Exec-managerial	Husband	White	Male	
...	...	...	...	...	...	
9760	Never-married	Sales	Own-child	White	Female	
9761	Divorced	Sales	Not-in-family	White	Female	
9762	Married-civ-spouse	Handlers-cleaners	Husband	White	Male	
9763	Married-civ-spouse	Other-service	Wife	White	Female	
9764	Never-married	Adm-clerical	Unmarried	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country	income	\
0	0	0	40	United-States	>50K	
1	0	0	45	United-States	<=50K.	
2	0	0	25	United-States	<=50K	
3	0	0	48	United-States	>50K.	
4	0	0	40	United-States	<=50K.	
...	...	...	...	...	...	
9760	0	0	30	United-States	<=50K	
9761	0	0	40	United-States	<=50K.	
9762	0	0	40	United-States	<=50K	
9763	3103	0	40	United-States	>50K.	
9764	0	0	40	United-States	<=50K	

	age_vs_hours
0	False
1	False
2	False
3	False
4	False
...	...
9760	False
9761	True
9762	False
9763	True
9764	False

[9765 rows x 16 columns]

##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####

This code receives and compares two numerical columns values (age and hour) across each row in the task\_dataset. Here, the compare\_age\_hours function takes a pandas DataFrame, task\_dataset, as its input, and for each row in the DataFrame, the function compares the value in the age column with the value in the hours-per-week column. It then creates a new column, age\_vs\_hours, where each entry is True if the individual's age is greater than the number of hours they work per week, and False otherwise.

The result\_task\_dataset then takes the compare\_age\_hours(task\_dataset) as its input and print it out with the print(result\_task\_dataset) function

(100 WORDS)

## Task 1.8: Write a function that returns the names of countries with maximum and minimum average ages in the task\_dataset (Hint: you can use the numpy module)

```
In [66]: ##### WRITE THE CODE IN THIS CELL#####
#Finding the countries with min and max average ages using Group by method
def find_extreme_average_age_countries():

    # Group by 'native-country' and calculate average age
    avg_age_by_country = task_dataset.groupby('native-country')['age'].mean()

    # Finding the country with maximum average age
    max_avg_age_country = avg_age_by_country.idxmax()

    # Find the country with minimum average age
    min_avg_age_country = avg_age_by_country.idxmin()

    return max_avg_age_country, min_avg_age_country

# Call the function and print the results
max_avg_country, min_avg_country = find_extreme_average_age_countries()
print("Country with maximum average age:", max_avg_country)
print("Country with minimum average age:", min_avg_country)
```

Country with maximum average age: Hungary

Country with minimum average age: Laos

##### WRITE EXPLANATIONS HERE (IF APPLICABLE) #####

The codes above seeks to analyze the task\_dataset to determine which countries have the maximum and minimum average age. respectively. The function calculates the average (mean) age for each country represented in the dataset by grouping the data by native-country and then applying the mean function to the age column. This results in a series object, avg\_age\_by\_country, indexed by country with average ages as values.



The country with the maximum average age and minimum average age are then identified using the `idxmax()` and `idxmin()` method on the `avg_age_by_country`. This method returns the index (country name) associated with the maximum and minimum value respectively